

Using the Cray Programming Environment

Alistair Hart

Cray Exascale Research Initiative Europe





Purpose of this talk

- **How to compile and run jobs**
 - If you have already used tödi, you may know some of this already
- **The practical exercises automate a lot of this**
 - The practicals are about learning OpenACC,
 - not remembering system-specific details
 - The Makefiles, jobscripts etc. can be used as templates for your projects



The system

- You are using a Cray system called "tödi"
 - it is a Cray XK7 system
 - you log in and compile on a front end node
 - these nodes have no GPU, so you can't run jobs here
 - jobs run on the compute nodes
 - each node has one AMD Interlagos CPU and one Nvidia GPU
 - 268 nodes contain Nvidia Kepler K20x GPUs
 - (4 nodes contain older Nvidia Fermi X2090 GPUs)
 - you run the jobs by submitting a jobscript to the SLURM batch system
 - compute jobs can't be run from the front end command line
 - there are two filesystems
 - home directories; yours is **\$HOME**
 - the lustre filesystem; your directory is **/scratch/todi/\$USER**
 - you should submit jobs from a directory on the lustre filesystem
 - home directories are backed up; the lustre filesystem is not and old files are periodically purged

Getting started

- **Cray uses a linux-based environment on the login nodes**
 - You will have a bash login shell by default
 - All the usual linux commands are available
 - Software versions are loaded and unloaded using the Gnu module command (see **man module**)
 - To see which modules are currently loaded, type: **module list**
 - To see which modules are available, type: **module avail**
 - You can wildcard the end of the names, e.g.: **module avail PrgEnv***
 - For more complicated grepping, you need to redirect stderr to stdout, e.g.
 - **module avail 2>&1 | grep "Env"**
 - You load a new module by typing: **module load <module name>**
 - Some modules (e.g. different compiler versions) conflict, so you should first **"module unload"** the old version (or use **"module swap"**)



Programming Environments

- **A number of different compilers are supported**

- You select these by loading a Programming Environment module
 - **PrgEnv-cray** for CCE (the default)
 - **PrgEnv-pgi** for PGI
 - **PrgEnv-gnu** for gcc, gfortran
- Once one of these is loaded, you can then select a compiler version
 - CCE: **module avail cce**
 - PGI: **module avail pgi**
 - Gnu: **module avail gcc**
- Swap to the most up to date version in each case
 - e.g. **"module avail cce"** to see the versions available
 - then **"module swap cce cce/<whatever>"**
- For any GPU programming (CUDA, OpenCL, OpenACC...)
 - make sure you always: **"module load craype-accel-nvidia35"**
 - it is not loaded by default



Using the compilers

- **You use the compilers via wrapper functions**

- `ftn` for Fortran; `cc` for C; `CC` for C++
- it doesn't matter which PrgEnv is loaded (same wrapper name)
- the wrappers add optimisation options, architecture-specific stuff and all the important library paths
 - make sure module `xtpe-interlagos` is loaded so these are correct
 - in many cases, you don't need any other compiler options
 - if you really want unoptimised code, you must use option `-O0`

- **Further information**

- man pages for the wrapper commands give you general information
- For more detail see the compiler-specific man pages
 - CCE: `crayftn`, `craycc`, `crayCC`
 - PGI: `pgfortran`, `pgcc`
 - GNU: `gfortran`, `gcc`
- You will need the appropriate PrgEnv module loaded to see these



Some Cray Compilation Environment basics

- **CCE-specific features:**

- Optimisation: **-O2** is the default and you should usually use this
 - **-O3** activates more aggressive options; could be faster or slower
- OpenMP: is supported by default.
 - if you don't want it, use **-hnoomp** compiler flag
- CCE only gives minimal information to stderr when compiling
 - to see more information, you should request a compiler listing file
 - flag **-hlist=a** for **ftn** and **cc**
 - writes a file with extension **.lst**
 - contains annotated source listing, followed by explanatory messages
 - each message is tagged with an identifier, e.g.: **ftn-6430**
 - to get more information on this, type: **explain <identifier>**
- For a full description of the Cray compilers, see the reference manuals at <http://docs.cray.com>.



Submitting jobs and the lustre filesystem

- **You should submit jobs from the lustre filesystem**
 - you can compile there as well if you wish
 - Create a unique directory for yourself: `mkdir -p /scratch/todi/$USER`
 - and subdirectories if you want
- **To submit a job, create a SLURM jobscript**
 - there is a skeleton script provided as part of the tutorial materials
 - just rename the executable
 - note that command `aprun` is used within the jobscript to run the executable
 - submit the job using command: `sbatch <jobscript name>`
 - other options are specified in the jobscript header
 - a job number is returned
 - to view the queued and running jobs: `squeue`
 - to see just your jobs: `squeue -u $USER`
 - to stop a queued or running job: `scancel <job number>`



A sample SLURM script

- **SLURM parameters:**

- **nodes** is the total number of nodes required
- **ntasks** is the number of MPI ranks
- **ntasks-per-node** is the number of ranks per node
 - Could be up to 32 for XE6
 - Usually set to 1 for XK7 GPU jobs (unless you are using CUDA_PROXY)
- **cpus-per-task** is number of threads per rank (usually 1 for GPU jobs)
 - $\text{ntasks-per-node} * \text{cpus-per-task} \leq 32$ for XE6
- See "**man sbatch**" for more details

- **Within the script:**

- you must manually set **OMP_NUM_THREADS** to **cpus-per-task** value
 - if you are using OpenMP
- run job using **aprun** command
 - "**aprun -B <executable>**" is a good shortcut
 - avoids need to specify "**-n**", "**-N**", "**-d**" options
 - **-n** should match **ntasks**
 - **-N** should match **ntasks-per-node**
 - **-d** should match **cpus-per-task**



A sample script

```
#!/bin/bash
#SBATCH --job-name="myjob"
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:01:00
#SBATCH --output="myjob.log"

# If you are using OpenMP, set the number of threads here
#export OMP_NUM_THREADS=<whatever>

# Now run the job with aprun
#   aprun -B uses the same parameters (-N -n -d) as specified
#   in PBS

# As an example, we'll run the PGI pgaccelinfo program to give
# information about the GPU. To access this, we need to load
# PrgEnv-pgi:
module swap PrgEnv-cray PrgEnv-pgi
# In general, you don't need to do this.

# Run the command:
aprun -B pgaccelinfo
```



For information: Compiling CUDA

- **Compilation:**

- `module load craype-accel-nvidia35`
- Main CPU code compiled with PrgEnv "`cc`" wrapper
 - either `PrgEnv-gnu` for gcc; or `PrgEnv-cray` for craycc
- GPU CUDA-C kernels compiled with `nvcc`
 - `nvcc -O3 -arch=sm_35`
- PrgEnv "`cc`" wrapper used for linking
 - Only GPU flag needed: `-lcudart`
 - e.g. no CUDA `-L` flags needed (added in `cc` wrapper)



For information: Compiling OpenCL

- **Compilation:**

- `module load craype-accel-nvidia35`
- Main CPU code compiled with PrgEnv "`cc`" wrapper
 - either `PrgEnv-gnu` for gcc; or `PrgEnv-cray` for craycc
- GPU OpenCL kernels compiled with `nvcc`
- PrgEnv "`cc`" wrapper used for linking
 - Only GPU flag needed: `-lOpenCL`

- **Alternatively:**

- Use `PrgEnv-gnu` for all compilation
 - still need `-lOpenCL` at linktime