



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

OpenACC Tutorial: Useful Strategies for Troubleshooting

Lugano, February 18th 2013

Jeff Poznanovic

Swiss National Supercomputing Center

poznanovic@cscs.ch



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

**With OpenACC, the
user gives much of the
control to the compiler
and runtime...**



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

Some Possible Issues

- **Code not running on the accelerator**
- **Unexpectedly slow execution**
- **Wrong answers**
- **Data not “present” on the accelerator**
- **Inconsistent host-device data**
 - Missing data clause
 - Missing update directive
- **And many types of CUDA API errors**
 - `CUDA_ERROR_LAUNCH_FAILED`
- ...



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

**Sorry, there's no
PRINTF() within
accelerator regions...**



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

So...

How do we

troubleshoot?



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

Overview of Troubleshooting Strategies

- 1. Check if an equivalent OpenMP version works on the host**
- 2. Look at the compiler messages for warnings**
- 3. Investigate the runtime output**
- 4. Compare with another OpenACC compiler/runtime**
- 5. Use Nvidia's tools**
- 6. Try a debugger**

This lecture's goal is to give you a toolbox of techniques in order to understand what the OpenACC compiler and runtime are doing.



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

1. Check the Multi-Core CPU Implementation

- **A good strategy: start with OpenMP and then move to OpenACC**
 - Debugging tools are more plentiful and robust on the host
 - Worry about host-device data movement at a later step
 - This allows you to focus on correct parallelization of your kernels
- **Gain confidence in your OpenMP CPU version first, and then the porting process to OpenACC will be much easier**

OpenMP

```
290. M-----< !$OMP PARALLEL SHARED (a,p,...) &
291. M          !$OMP PRIVATE (s0,...)
292. M 2-----< do loop=1,nn
[... ]
298. M 2
299. M 2 m-----< do k=2,kmax-1
300. M 2 m 4-----< do j=2,jmax-1
301. M 2 m 4 Vr3--< do i=2,imax-1
302. M 2 m 4 Vr3      s0=a(I,J,K,1)*p(I+1,J,K) &
303. M 2 m 4 Vr3      +a(I,J,K,2)*p(I,J+1,K) &
304. M 2 m 4 Vr3      +a(I,J,K,3)*p(I,J,K+1) &
```



OpenACC

```
290.          !$ACC DATA COPYIN(a,p,...)
291. 1-----< do loop=1,nn
[... ]
294. 1 G-----<
295. 1 G g-----< !$ACC KERNELS LOOP PRIVATE(s0,...)
296. 1 G g b-----< do k=2,kmax-1
297. 1 G g b gb--< do j=2,jmax-1
298. 1 G g b gb      do i=2,imax-1
299. 1 G g b gb      s0=a(I,J,K,1)*p(I+1,J,K) &
300. 1 G g b gb      +a(I,J,K,2)*p(I,J+1,K) &
                +a(I,J,K,3)*p(I,J,K+1) &
```



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

2. Compiler Messages: Cray

- **Usage:** `ftn -h list=m... / cc -h list=m ...`

Then, look at the loopmark file (*.lst)

POSITIVE:

```
224.          #pragma acc parallel pcopy(qleft[0:Hnvar*Hstep*Hnxyt] ...)
225.          #pragma acc loop gang collapse(2)
226.  + GC-----<      for (invar = IP + 1; invar < Hnvar; invar++) {
227.  + GC g----<        for (s = 0; s < slices; s++) {
228.      GC g            #pragma acc loop vector
229.      GC g g--<        for (i = 0; i < narray; i++) {
230.      GC g g          [...]
```

```
CC-6405 CC: ACCEL File = riemann.c, Line = 226
A region starting at line 226 was placed on the accelerator.
```

```
CC-6418 CC: ACCEL File = riemann.c, Line = 226
If not already present: allocate memory and copy user shaped variable "qleft" to
accelerator, free at line 234 (acc_copyin).
```

Use pre-allocated device data

```
CC-6060 CC: SCALAR File = riemann.c, Line = 226
A loop nest was collapsed according to user directive.
```

```
CC-6430 CC: ACCEL File = riemann.c, Line = 227
A loop was partitioned across the thread blocks.
```

Successful partitioning of loops
and mapping to the hardware

```
CC-6430 CC: ACCEL File = riemann.c, Line = 229
A loop was partitioned across the 128 threads within a threadblock.
```

NEGATIVE:

```
=====
Examples of information you should pay attention to:
  Loop was not vectorized because a recurrence was found
  A loop starting at line 111 will be serially executed.
  A loop starting at line 222 will be redundantly executed.
```

IMPORTANT:

Messages can hint at issues
with kernel generation,
performance, and correctness



2. Compiler Messages: PGI

- **Usage:** `ftn -Minfo=accel ... / cc -Minfo=accel ...`
- **Misc tips:** http://www.pgroup.com/resources/openacc_tips_fortran.htm
<http://www.pgroup.com/lit/articles/insider/v1n2a1.htm>

POSITIVE:

```
pgcc -acc -ta=nvidia -Minfo=accel [...] -c riemann.c
Dmemset:
  37, Generating present_or_copyout(t[0:nbr])
      Accelerator kernel generated
  38, #pragma acc loop gang, vector(256) /* blockIdx.x threadIdx.x */
  37, Generating NVIDIA code
      Generating compute capability 3.0 binary
riemann:
  78, Generating present(sgnm[0:Hnxyt*Hstep])
      Generating present(qgdnv[0:Hnxyt*(Hstep*Hnvar)])
      Generating present(qright[0:Hnxyt*(Hstep*Hnvar)])
      Generating present(qleft[0:Hnxyt*(Hstep*Hnvar)])
      Accelerator kernel generated
  80, #pragma acc loop gang /* blockIdx.x */
  82, #pragma acc loop vector(256) /* threadIdx.x */
  78, Generating NVIDIA code
      Generating compute capability 3.0 binary
```

Use pre-allocated device data

Successful partitioning of loops and mapping to the hardware

NEGATIVE:

```
=====
Examples of information you should pay attention to:
  Loop carried scalar dependence for 'x'
  Parallelization would require privatization of array 'y[0:n-1]'
  Accelerator restriction: scalar variable live-out from loop: z
  Accelerator restriction: unsupported operation
  Accelerator restriction: feature XYZ not supported
```

IMPORTANT:
Messages can hint at issues with kernel generation, performance, and correctness



2. Compiler Messages: CAPS

- **Usage:** `hmpp --debug cc/ftn ...`
- **Example troubleshooting scenario:** why is my kernel code so slow?

SLOW:

```
!$ACC PARALLEL LOOP
do k=2,kmax-1
  do j=2,jmax-1
    do i=2,imax-1
      [...]
```

```
pjeffrey@todil> hmpp --debug ftn -o test1 test1.f90
hmppcg: [Message DPL0099] test1.f90:308: Loop 'k' was shared among gangs(16)
hmpp: [Info] Generated codelet filename is "hmpp_acc_region__k4ehwytw__cuda.hmf.cu".
pjeffrey@todil> aprun ./test1
MFLOPS: 148.94
```

No message for
vector parallelism

Unexpectedly slow

FAST:

```
!$ACC PARALLEL LOOP GANG
do k=2,kmax-1
  do j=2,jmax-1
    !$ACC LOOP VECTOR
    do i=2,imax-1
      [...]
```

```
pjeffrey@todil> hmpp --debug ftn -o test2 test2.f90
hmppcg: [Message DPL0099] test2.f90:311: Loop 'i' was vectorized(128)
hmppcg: [Message DPL0099] test2.f90:308: Loop 'k' was shared among gangs(16)
hmpp: [Info] Generated codelet filename is "hmpp_acc_region__k4ehwytw__cuda.hmf.cu".
pjeffrey@todil> aprun ./test2
MFLOPS: 5136.80
```

Messages for
gang+vector
parallelism

Good performance



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

3. Runtime Output: Cray

• Usage: **export CRAY_ACC_DEBUG=[0-3]**

```
[...]
ACC: Start transfer 1 items from data_locality.f90:47
ACC: flags:
ACC:
ACC: Trans 1
ACC: Simple transfer of 'array' (4000 bytes)
ACC: host ptr 1000003b000
ACC: acc ptr 0
ACC: flags: ACQ_PRESENT REG_PRESENT
ACC: host region 1000003b000 to 1000003bfa0 found in present table index 0 (ref count 2)
ACC: memory found in present table (b00200000, base b00200000)
ACC: new acc ptr b00200000
ACC:
ACC: End transfer (to acc 0 bytes, to host 0 bytes)
ACC:
ACC: Start kernel my_sub_$ck_L48_3 async(auto) from data_locality.f90:48
ACC: flags: CACHE_MOD CACHE_FUNC AUTO_ASYNC FLEX_BLOCKS
ACC: mod cache: 0x6039c0
ACC: kernel cache: 0x603980
ACC: async info: 0x2aaabfa6a460
ACC: arguments: NVIDIA argument info
ACC: param size: 12
ACC: param pointer: 0x7fffff9dc0
ACC: blocks: 10
ACC: threads: 128
ACC: event id: 0
ACC: using cached module
ACC: getting function
ACC: stats threads=1024 threadblocks per sm=8 shared=0 total shared=0
ACC: prefer L1 cache
ACC: kernel information
ACC: num registers : 25
ACC: max theads per block : 1024
ACC: shared size : 0 bytes
ACC: const size : 0 bytes
ACC: local size : 0 bytes
ACC:
ACC: launching kernel new
ACC: caching function
ACC: End kernel
[...]
```

Some explanation:

Found 'array' on the device
and use its device
address, ...

Launch the kernel with 10
blocks and 128 threads/
block.

The kernel uses 25
registers, ...

**CSCS**Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

3. Runtime Output: PGI

- **Usage:** **export PGI_ACC_DEBUG=[0,1] (more)**
export ACC_NOTIFY=[0-3] (less)

Some explanation:

```
[...]
pgi_acc_dataon(devptr=0x436fd5,hostptr=0x73c1a0,offset=0,0,0,stride=1,10,100,
              size=10x10x10,extent=10x10x10,eltsize=4,lineno=47,name=array,
              flags=0x201=sync+present)
mapped host:0x73c1a0 dev:0xb00200000 offset:0 (host:0x73c1a0 dev:0xb00200000
size:4000 offset:0 data[host:0x73c1a0 dev:0xb00200000 size:4000]
line:18 name:my_array)
__pgi_cu_init( file=data_locality.f90, function=my_sub, line=48, startline=41, endline=58 )
__pgi_cu_module3( lineno=48 )
__pgi_cu_module3 module loaded at 0xa84ca0
__pgi_cu_module_function( name=0x6ede1a=my_sub_48_gpu, lineno=48, argname=(nil)=,
argsize=56, varname=(nil)=, varsize=0, SWcachesize=0 )
Function handle is 0xa2cb50
__pgi_cu_alloc(size=40,lineno=48,name=k)
__pgi_cu_alloc(40) returns 0xb00300000 (address=0x7fffffff048)
__pgi_cu_alloc(size=40,lineno=48,name=j)
__pgi_cu_alloc(40) returns 0xb00300200 (address=0x7fffffff050)
__pgi_cu_alloc(size=40,lineno=48,name=i)
__pgi_cu_alloc(40) returns 0xb00300400 (address=0x7fffffff058)
__pgi_cu_launch_a(func=0xa2cb50, grid=10x1x1, block=256x1x1, lineno=48)
__pgi_cu_launch_a(func=0xa2cb50, params=0x7fffffff000, bytes=56, sharedbytes=0)
First arguments are:
      10      10      10      0  2097152      11      10      10
3146752      11      10      111      10      100

0x0000000a 0x0000000a 0x0000000a 0x00000000 0x00200000 0x0000000b
0x00300400 0x0000000b 0x0000000a 0x0000000f 0x0000000a 0x00000064
[...]
```

Found 'array' on the device
and use its device
address, ...

Some initialization

Allocate some local storage

Launch the kernel with 10
blocks and 256 threads/
block, ...



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

3. Runtime Output: CAPS

• **Usage:** `export HMPVRT_LOG_LEVEL=[off,info,all,...]`

```
[...]
[ 0.486952] ( 0) INFO : Enter data (queue=none, location=data_locality.f90:47)
[ 0.486981] ( 0) INFO : Enter parallel (queue=none, location=data_locality.f90:48)
[ 0.487009] ( 0) DEBUG: Grouplet(0x69cce0)::Grouplet(file_name='hmpp_acc_region__vx8sok80__cuda.hmf')
[ 0.489138] ( 0) DEBUG: Grouplet(0x69cce0)::setTarget(target=cuda)
[ 0.489171] ( 0) DEBUG: Grouplet(0x69cce0)::addCodelet(codelet_name='hmpp_acc_region__vx8sok80_',
signature_string='prototype hmpp_acc_region__vx8sok80_(n: s32, array: inout ^cudaglob s32)')
[ 0.489235] ( 0) DEBUG: Codelet(0x10000091590)::Codelet(grouplet=Grouplet(0x69cce0),
name="hmpp_acc_region__vx8sok80_", function=0x2aaac81f8a74,
signature=Signature(0x100000410e0))
[ 0.489282] ( 0) DEBUG: Grouplet(0x69cce0)::getCodeletByName(codelet_name='hmpp_acc_region__vx8sok80_')
[ 0.489306] ( 0) INFO : Call hmpp_acc_region__vx8sok80_ (queue=none, location=data_locality.f90:48)
[ 0.489332] ( 0) DEBUG: Codelet(0x10000091590)::call(device=Device(0x6b89e0), arguments=...<2>,
queue=(nil))
[ 0.489392] ( 0) DEBUG: CUDADevice(0x6b89e0)::allocate(memory_space=cudashared, size=4)
[ 0.489420] ( 0) DEBUG: CUDADevice(0x6b89e0)::allocate(memory_space=cudashared, size=4)
[ 0.489447] ( 0) DEBUG: CUDADevice(0x6b89e0)::allocate(memory_space=cudashared, size=4)
[ 0.489502] ( 0) DEBUG: CUDADevice(0x6b89e0)::launchGrid(grid=CUDAGrid(0x6d98d0),
call=CUDAGridCall(0x7fffff90c0), queue=(nil))
[ 0.489541] ( 0) DEBUG: cuModuleLoadDataEx(module=..., image=0x2aaac81f8ba0, numOptions=5,
options=..., optionValues=...)
[ 0.489740] ( 0) DEBUG: -> *module=CUmodule(0x10000068700)
[ 0.489769] ( 0) DEBUG: cuModuleGetFunction(hfunc=..., hmod=CUmodule(0x10000068700),
name='hmpp_acc_region__vx8sok80__parallel_region_1')
[ 0.489802] ( 0) DEBUG: -> *hfunc=CUfunction(0x100004b3180)
[ 0.489827] ( 0) DEBUG: cuLaunchKernel(f=CUfunction(0x100004b3180), gridDimX=16, gridDimY=1, gridDimZ=1,
blockDimX=32, blockDimY=8, blockDimZ=1, sharedMemBytes=0,
hStream=CUstream((nil)), kernelParams=..., extra=...)
[ 0.489909] ( 0) DEBUG: cuCtxSynchronize()
[ 0.489965] ( 0) DEBUG: CUDADevice(0x6b89e0)::free(memory_space=cudashared, address=0x100000914d0)
[ 0.489999] ( 0) DEBUG: CUDADevice(0x6b89e0)::free(memory_space=cudashared, address=0x100000917a0)
[ 0.490026] ( 0) DEBUG: CUDADevice(0x6b89e0)::free(memory_space=cudashared, address=0x10000091410)
[ 0.490054] ( 0) INFO : Leave parallel (queue=none, location=data_locality.f90:48)
[ 0.490081] ( 0) INFO : Leave data (queue=none, location=data_locality.f90:47)
[...]
```

Some explanation:

Hit the data and
parallel regions

Some initialization

Allocate some
shared storage

Launch the kernel
with 16 blocks and
256 threads/
block, ...

Free some storage

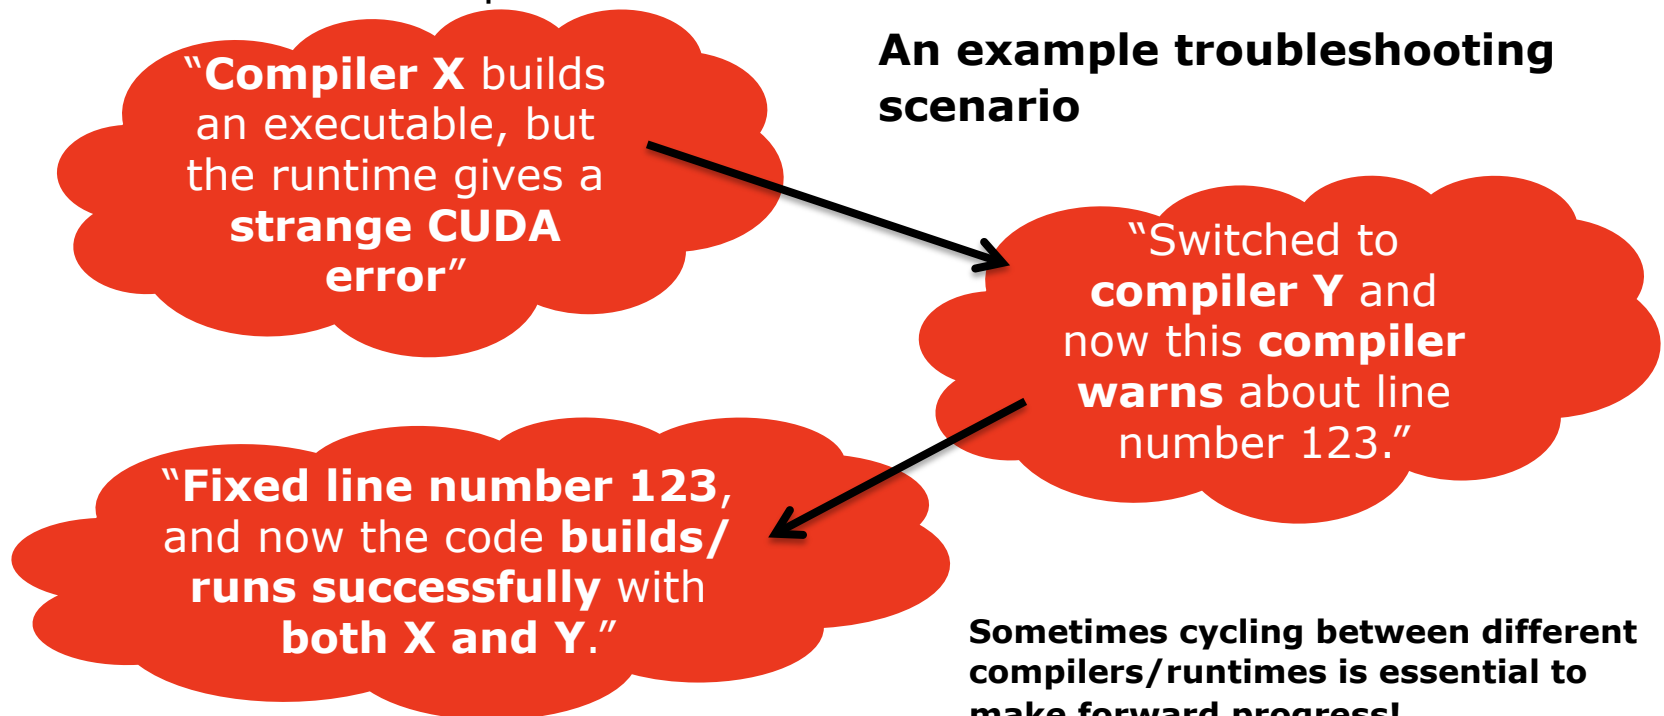


CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

4. Comparative Debugging

- **Extremely useful: comparing the compiler messages and runtime output from CAPS, Cray and PGI**
 - A huge amount of information is available at compile-/run-time
 - Each compiler vendor has its own strengths and weaknesses
 - Each vendor provides a different set of information





CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

5. Nvidia Tools for OpenACC: CUDA Memcheck

- **Cuda-memcheck is included with all CUDA installations**
- **Example troubleshooting scenario**
 - CUDA_ERROR_LAUNCH_FAILED runtime error

The following output gives the function and line number of the accelerator region that causes an out-of-bound memory access:

```
aprun -n1 cuda-memcheck ./a.out
```

```
===== CUDA-MEMCHECK
```

```
100000 iterations completed
```

```
===== Invalid __global__ write of size 4
```

```
===== at 0x000000f8 in main_$ck_L21_1
```

```
===== by thread (32,0,0) in block (781,0,0)
```

```
===== Address 0x2002c3680 is out of bounds
```

```
=====
```

```
===== ERROR SUMMARY: 1 error
```

Function

Line number

Error

More details on this tool will come in the CUDA section

Link: <https://developer.nvidia.com/cuda-memcheck>

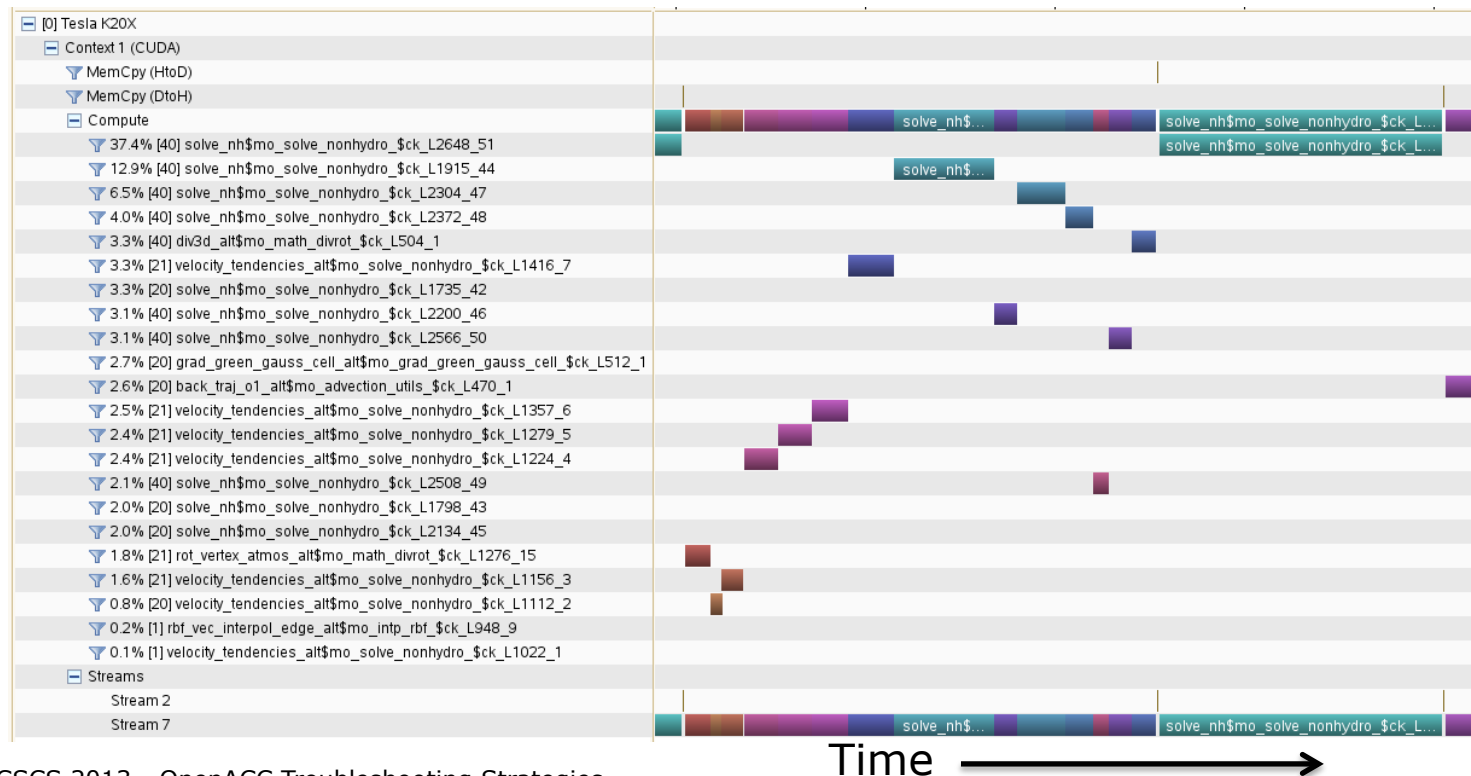


CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

5. Nvidia Tools for OpenACC: Timeline GUIs

- **Timeline GUIs can be useful for visualizing host-device events**
- **Example troubleshooting scenarios**
 - Visualizing the ordering of kernels and data transfers
 - Verifying any expected asynchronous behavior





CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

6. Using the Allinea DDT Debugger for OpenACC

- **Requires DWARF support (compiler-generated debugger info)**
 - PGI: work in progress, debug on host with “-ta=host”
 - CAPS: uses nvcc backend, matches with orig src code
 - Cray: host and device support
- **DDT’s current status**
 - Support for CUDA and OpenACC; some ongoing issues
 - DDT 4.0 includes important fixes from Nvidia (soon to be released)
- **Capabilities (CUDA and OpenACC)**
 - Set breakpoints in host and device code
 - Step through kernels with various blockIdx and threadIdx configs
 - Inspect variables and arrays on the host and device
- **Usage (with Cray CCE)**
 - module load ddt
 - module load craype-accel-nvidia35
 - ftn **-g** -h acc [...] / cc **-g** -h acc [...] / nvcc **-g -G** [...]
 - ddt ./a.out
- **Notes**
 - Non-MPI codes may need to add a dummy MPI_Init() call
 - You may need to decrease optimization to preserve src line info



6. Using the Allinea DDT Debugger for OpenACC

The screenshot shows the Allinea DDT v3.2.1-27702 debugger interface. The main window displays the source code of a program, with line 311 highlighted. The left sidebar shows the 'Threads' pane with thread 1 selected. The right sidebar shows the 'Locals' pane with variables like 'a', 'b', 'bnd', 'i', 'imax', 'j', 'jmax', 'k', 'kmax', 'omega', 'p', 's0', 'ss', 'wgosa', and 'wrk1'. The bottom pane shows the 'Stacks' pane with the current stack frame.

Execution control (indicated by a red arrow pointing to the execution controls at the top left)

Breakpoints (indicated by a red arrow pointing to the breakpoint icon on the left side of the source code window)

Current accelerator region (indicated by a red arrow pointing to the 'jacob_i_\$ck_l306_9' entry in the 'Stacks' pane)

blockIdx/threadIdx (indicated by a red arrow pointing to the 'Block' and 'Thread' fields in the 'CUDAThreads' pane)

Local variables/arrays (indicated by a red arrow pointing to the 'Locals' pane on the right)



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

**Give these techniques a
try, and let us know
about your experiences**

help@cscs.ch

Thanks!