

Porting real applications

Alistair Hart

Cray Exascale Research Initiative Europe





Adding OpenACC to a Larger Code

- **Adding OpenACC to a real code is not trivial work...**
 - Are parts of the program suitable for an accelerator?
 - Where do we start?
 - What do we do next?
- **We'll go through the exercise for an example code**
 - Running on a Cray XK7 (AMD Interlagos and Nvidia Kepler K20x)
 - Using Cray compiler and Cray performance analysis tools



The Code

- **NAS Parallel Benchmarks MG (MultiGrid) code**
 - Shorter than typical application
 - but structure of code is very similar
 - This example concentrates on the serial version
 - We also have parallel versions ported to OpenACC
 - The serial versions have OpenMP directives, but we do not use them during this exercise
 - Downloading it:
 - Fortran version: <http://www.nas.nasa.gov/publications/npb.html>.
 - 1445 lines, of which 267 blank
 - C version: <http://www.hpcs.cs.tsukuba.ac.jp/omni-openmp/download/download-benchmarks.html>.
 - 1292 lines, of which 206 blank



Building and Running MG

● Build:

- `make MG [CLASS=<CLASS>] [<OPTIONS>]`
 - `CLASS` is the problem size. "B" is the default.
- Top-level `Makefile` passes options to `MG/Makefile`
 - this uses `config/make.def` for compiler-specific options

● Run:

- Three important lines of output
 - Fortran
 - `L2 Norm is 0.1800564401355E-05`
 - `Mop/s total = 1623.04`
 - `Verification = SUCCESSFUL`
 - C output same, but baseline performance differs
 - `Mop/s total = 1320.26`
- Always check:
 - L2 Norm should not be a NaN
 - Verification should be successful



Where Do We Start?

- Profile MG on the CPU

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function
100.0%	12.069520	--	--	1630.0	Total

100.0%	12.069417	--	--	1230.0	USER

54.9%	6.620529	--	--	161.0	resid_
25.3%	3.057070	--	--	160.0	psinv_
9.5%	1.148982	--	--	140.0	rprj3_
8.1%	0.983395	--	--	140.0	interp_
1.3%	0.153775	--	--	461.0	comm3_
=====					

- Four routines dominate the runtime
 - More than half the time is spent in **resid**
 - There are other routines executing for less than 1% of the total time
 - These might be important for the OpenACC port



Understand Flow of the Application

Table 1: Function Calltree View

Time%	Time	Calls	Calltree
100.0%	12.069520	1630.0	Total

100.0%	12.069417	1230.0	mg_

72.3%	8.724588	1180.0	mg3p_

3 28.3%	3.416675	280.0	resid_
3 25.8%	3.108020	320.0	psinv_
3 9.6%	1.160157	280.0	rprj3_
3 8.1%	0.983395	140.0	interp_
=====			
27.3%	3.295504	42.0	resid_
=====			

- **mg calls:**
 - mg3p (which then calls resid, psinv, rprj3, interp)
 - resid also called directly from mg

Get Work Estimates for Loops

Table 2: Loop Stats by Function (from -hprofile_generate)

Loop Incl Time Total	Loop Hit	Loop Trips Avg	Loop Trips Min	Loop Trips Max	Function=/.LOOP[.]
6.830878	161	96.497	4	256	resid_.LOOP.1.li.634
6.830032	15536	201.067	4	256	resid_.LOOP.2.li.635
4.033780	3123776	237.548	6	258	resid_.LOOP.3.li.636
2.607888	3123776	235.548	4	256	resid_.LOOP.4.li.642

- **Loop-level profiling is more useful now**

- Which loopnests (rather than just routines) took most time?
- How many iterations did this loopnest have?

- **Here are the lines relating to resid**

- Loops starting at 636 and 642 are nested inside loops at line 634, 635
 - See how the Loop Hit numbers multiply up
 - See how inclusive times for 636 and 642 add to give that for 635
 - Inclusive times for 634, 635 same: perfectly nested loops



Add First OpenACC Kernel

- Clearly we should start with **resid**
 - Fortran:
 - `!$acc parallel loop vector_length(NTHREADS)`
`!$acc& private(u1,u2) copyin(u,v,a) copyout(r)`
 - C:
 - `#pragma acc parallel loop vector_length(NTHREADS) \`
`private(u1,u2) copyin(u[0:n1*n2*n3],v[0:n1*n2*n3],a[0:4]) \`
`copyout(r[0:n1*n2*n3])`
 - Data movement sizes explicit to avoid "unshaped pointer" errors
 - Because we are dynamically allocating memory



Resulting MG Performance?

- Running with and without OpenACC kernel:

	Original (Mop/s)	1 kernel (Mop/s)
Fortran	1623.04	1541.42
C	1320.26	1409.48

- So the code is actually slower... **Why?**



Enable Cray Runtime Commentary

- export `CRAY_ACC_DEBUG=2`
- for every call to `resid`:

```
ACC: Start transfer 6 items from mg_v03.f:615
ACC:      allocate, copy to acc 'a' (32 bytes)
ACC:      allocate 'r' (137388096 bytes)
ACC:      allocate, copy to acc 'u' (137388096 bytes)
ACC:      allocate, copy to acc 'v' (137388096 bytes)
ACC:      allocate <internal> (530432 bytes)
ACC:      allocate <internal> (530432 bytes)
ACC: End transfer (to acc 274776224 bytes, to host 0 bytes)
ACC: Execute kernel resid_$ck_L615_1 blocks:256 threads:128 async(auto) from mg_v03.f:615
ACC: Wait async(auto) from mg_v03.f:639
ACC: Start transfer 6 items from mg_v03.f:639
ACC:      free 'a' (32 bytes)
ACC:      copy to host, free 'r' (137388096 bytes)
ACC:      free 'u' (137388096 bytes)
ACC:      free 'v' (137388096 bytes)
ACC:      free <internal> (0 bytes)
ACC:      free <internal> (0 bytes)
ACC: End transfer (to acc 0 bytes, to host 137388096 bytes)
```

- **Certainly a lot of data was moved**
 - Commentary tells us which arrays, at which line and how much data



Or Use Nvidia Compute Profiler

- **export COMPUTE_PROFILE=1**

- Analyses PTX (from OpenACC or from CUDA)
- Very useful if mixing OpenACC with CUDA code

```
method=[ memcpyHtoD ] gputime=[ 1.088 ] cputime=[ 42.000 ]
method=[ memcpyHtoD ] gputime=[ 52236.543 ] cputime=[ 52513.000 ]
method=[ memcpyHtoD ] gputime=[ 52153.281 ] cputime=[ 52402.000 ]
method=[ resid_$ck_L615_1 ] gputime=[ 15063.424 ] cputime=[ 21.000 ] occupancy=[ 0.333 ]
method=[ memcpyDtoH ] gputime=[ 281508.594 ] cputime=[ 283700.000 ]
```

- **Data transfers obvious, taking most time**

Or Use CrayPAT for a Profile by Function

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function
100.0%	16.409900	--	--	1252.0	Total
100.0%	16.409731	--	--	851.0	USER
51.2%	8.403343	--	--	1.0	mg_
34.3%	5.622111	--	--	170.0	resid_.ACC_COPY@li.615
11.8%	1.936478	--	--	170.0	resid_.ACC_COPY@li.639
2.7%	0.440894	--	--	170.0	resid_.ACC_SYNC_WAIT@li.639
0.0%	0.005727	--	--	170.0	resid_.ACC_KERNEL@li.615
0.0%	0.001178	--	--	170.0	resid_.ACC_REGION@li.615
0.0%	0.000170	--	--	401.0	ETC

- Provides aggregated report of data movements
 - names, sizes and frequencies of original arrays lost
- Shows asynchronous kernel launches
 - Notice ACC_KERNEL almost zero
 - SYNC_WAIT shows the compute time
 - could recompile with **-hacc_model=auto_async_none**

...And CrayPAT Accelerator Statistics

Table 2: Time and Bytes Transferred for Accelerator Regions

Host Time%	Host Time	Acc Time	Acc Copy In (MBytes)	Acc Copy Out (MBytes)	Events	Calltree
100.0%	8.007	7.962	12341	6171	850	Total
100.0%	8.007	7.962	12341	6171	850	mg_
50.0%	4.005	3.969	6314	3157	735	mg3p_
3						resid_
4						resid_.ACC_REGION@li.615
5	36.2%	2.898	2.877	6314	--	147 resid_.ACC_COPY@li.615
5	10.8%	0.867	0.860	--	3157	147 resid_.ACC_COPY@li.639
5	2.9%	0.235	--	--	--	147 resid_.ACC_SYNC_WAIT@li.639
5	0.1%	0.004	0.232	--	--	147 resid_.ACC_KERNEL@li.615
5	0.0%	0.001	--	--	--	147 resid_.ACC_REGION@li.615(exclusive)

- Host and accelerator times given separately

- ACC_KERNEL

- Acc Time is the compute time
 - Host Time is the time for the asynchronous launch
 - The Host "catches up" at the SYNC_WAIT



... And CrayPAT Summarized Trace

- pat_build -u mg.B.x

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function
100.0%	16.452925	--	--	265303.0	Total
100.0%	16.452760	--	--	264902.0	USER
34.2%	5.621172	--	--	170.0	resid_.ACC_COPY@li.615
19.4%	3.199216	--	--	168.0	psinv_
11.8%	1.940111	--	--	170.0	resid_.ACC_COPY@li.639
10.7%	1.764268	--	--	131072.0	vranlc_
7.4%	1.217534	--	--	147.0	rprj3_
6.3%	1.033920	--	--	147.0	interp_
4.3%	0.709337	--	--	151.0	zero3_
2.7%	0.441237	--	--	170.0	resid_.ACC_SYNC_WAIT@li.639
1.5%	0.240856	--	--	2.0	zran3_
1.0%	0.170554	--	--	487.0	comm3_

- resid kernel no longer dominates the profile

- actual compute time is shown in SYNC_WAIT (Host) Time
- Its data copies are significant, however



More OpenACC Kernels

- Running with 4 accelerated kernels:

	Original (Mop/s)	1 kernel (Mop/s)	4 kernels (Mop/s)
Fortran	1623.04	1541.42	1274.48
C	1320.26	1409.48	991.42

- Even slower, and C particularly bad. **Why?**

Profile the Code Again

- Notice that spending most time in interp
- Next look at compiler listing:

```

727.      #pragma acc parallel loop private(z1,z2,z3) \
728.      copy(u[0:n1*n2*n3])
729.      copyin(z[0:mm1*mm2*mm3])
730.  gG-----<      for (i3 = 0; i3 < mm3-1; i3++) {
731.  1-----<          for (i2 = 0; i2 < mm2-1; i2++) {
732.  1 g-----<          for (i1 = 0; i1 < mm1; i1++) {
733.  1 g              i123 = i1 + mm1*i2 + mm12*i3;
734.  1 g              z1[i1] = z[i123+mm1] + z[i123];
735.  1 g              z2[i1] = z[i123+mm12] + z[i123];
736.  1 g              z3[i1] = z[i123+mm1+mm12] + z[i123+mm12] + z1[i1];
737.  1 g----->          }
738.  1 r4-----<      for (i1 = 0; i1 < mm1-1; i1++) {
739.  1 r4          i123 = i1 + mm1*i2 + mm12*i3;
740.  1 r4          j123 = 2*i1 + n1*(2*i2 + n2 * 2*i3);
741.  1 r4          u[j123] += z[i123];
742.  1 r4          u[j123+1] += 0.5*(z[i123+1]+z[i123]);
743.  1 r4----->      }

```

Loop
Accelerated

Loop
Partitioned

Loop Not
Partitioned

- Loop at line 738 not partitioned
 - Executed redundantly: every thread does every loop iteration



More OpenACC Kernels

- Insert directive above unpartitioned loop to help compiler:
 - `#pragma acc loop independent`
 - (if this didn't work, would use "`#pragma acc loop vector`" instead)

	Original (Mop/s)	1 kernel (Mop/s)	4 kernels (Mop/s)
Fortran	1623.04	1541.42	1274.48
C	1320.26	1409.48	1271.12

- Fortran and C performance now identical



Next Steps – Reduce Data Movement

- Need to introduce data regions higher up calltree
- For this, need all callee routines to be accelerated with OpenACC directives
- Use a profiler to map out the calltree to get list of routines
- First need to port some insignificant routines
 - norm2u3, zero3, comm3



Results for Accelerating up the Calltree

- Accelerated loops in norm2u3, zero3, comm3

	Original (Mop/s)	1 kernel (Mop/s)	4 kernels (Mop/s)	Calltree Routines (Mop/s)
Fortran	1623.04	1541.42	1274.48	886.02
C	1320.26	1409.48	1271.12	873.01

- Slower because even more data movement

- C still slightly down; poor scheduling needs acc loop independent

	Original (Mop/s)	1 kernel (Mop/s)	4 kernels (Mop/s)	Calltree Routines (Mop/s)
Fortran	1623.04	1541.42	1274.48	886.02
C	1320.26	1409.48	1271.12	885.27



Add Data Region

- **Now we put a data region in the main program**
 - Arrays u,v,r are declared **create**
 - We'll never use the host version of these
 - Arrays a,c are declared **copyin**
 - They're initialized on the host
- **Then in all the subprograms, we change clauses**
 - Replace **copy*** and **create** by **present**
 - Could replace by **present_or_***
 - If we know they should always be present, better to state this
 - Then mistakes become runtime errors rather than just wrong answers
 - We'd have to diagnose these by trawling the runtime commentary



Results with Data Region in Main

- At last, we are running faster (and correctly)!

	Original (Mop/s)	1 kernel (Mop/s)	4 kernels (Mop/s)	Calltree Routines (Mop/s)	Data Region (Mop/s)
Fortran	1623.04	1541.42	1274.48	886.02	23913.21
C	1320.26	1409.48	1271.12	885.27	23558.03



View Compiler Commentary Again

- **All array data transfers eliminated**

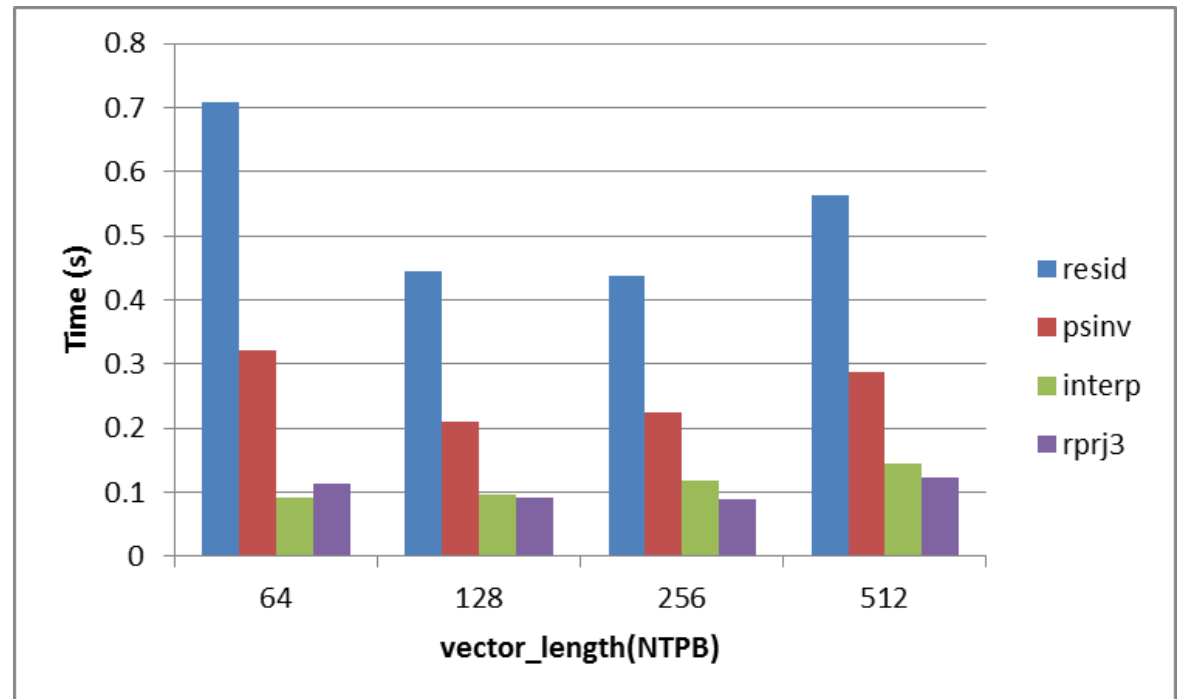
- Run with `CRAY_ACC_DEBUG=2` and catch STDERR in a file
- `grep "copy" <file> | sort | uniq`

```
ACC:      allocate, copy to acc 'a' (32 bytes)
ACC:      allocate, copy to acc 'c' (32 bytes)
ACC:      allocate, copy to acc 'jg' (320 bytes)
ACC:      allocate reusable, copy to acc <internal> (16 bytes)
ACC:      allocate reusable, copy to acc <internal> (4 bytes)
ACC:      copy to host, done reusable <internal> (16 bytes)
ACC:      reusable acquired, copy to acc <internal> (16 bytes)
```

- Arrays a, c, jg copied only at initialization
- Some internal transfers unavoidable

Performance Tuning Tips

- **Check the .lst loopmark file**
 - Are any kernels obviously badly scheduled?
 - No, we already checked that
- **Try varying vector_length from the default of 128**
 - Different values may suit different kernels
 - Effect small here: up to 5% per kernel, but only 2% overall:





Performance Tuning Tips

● Loop scheduling

- Collapse loops within loopnests
 - OpenACC schedules according to the loops in the nest
 - Collapsing loops can increase the tripcount
 - e.g. to allow more threads in a block
- Using the worker clause may help for imperfectly nested loops
- Block loops in a loopnest
 - This can improve cache usage (as on the CPU)
 - CCE-specific directives can do this, or try it manually

● More extreme:

- Avoid temporary arrays
 - Use private scalars, as these are more likely to go into registers
- Rewrite the most expensive kernels in CUDA and handtune
 - but remember you are competing against a whole compiler team



Other Performance Tuning Improvements

- **Avoiding temporary arrays in resid:**

- Mop/s total = 23999.08 ! Fortran
- Mop/s total = 23640.24 // C

- benefit v. small; was it really worth hacking the source?

- **Call an external CUDA version of resid**

- Mop/s total = 22351.51 ! Fortran
- Mop/s total = 21289.72 // C

- This was a naive kernel
 - (Even so, there may be a lesson in this)

- **Data movement was a far bigger optimisation**

- than any of our kernel improvements



Conclusions: How far did we get?

- **Significant speedup compared to single core:**

- Fortran: $1826.19 \rightarrow 23999.08 = 15x$
- C: $1320.26 \rightarrow 23640.24 = 18x$

- **The real comparison is to a full CPU or node**

- run the OpenMP version of the code
 - across 16 cores for an XK6 node (single AMD Interlagos)
 - Mop/s total = 9162.37 ! Fortran, Cray XK7 CPU, 16 threads
 - Mop/s total = 8638.68 // C, Cray XK7 CPU, 16 threads
 - across 32 cores for an XE6 node (dual AMD Interlagos)
 - Mop/s total = 15244.09 ! Fortran, Cray XE6, 32 threads
 - Mop/s total = 15120.86 // C, Cray XE6, 32 threads
- maybe MPI version would scale better, but our code here is scalar

- **Final speedup compared to full node:**

- XK7 CPU node (16 cores): **2.6x**
- XE6 CPU node (32 cores): **1.5x**
 - This is for Fortran, C looks slightly better



Conclusions: How much further could we get?

- **So we are 2.6x or 1.5x faster, node-for-node**
 - How much faster could we get (speeds and feeds)
 - Flops and mem. b/w around 5-10x faster than single AMD Interlagos
- **So why were we not faster?**
 - MultiGrid application cycles through grid sizes
 - sometimes the grid is really small: 4x4x4
 - CrayPAT loop profiling showed us that
 - Small grid sizes will never schedule well on the GPU
 - consider checking grid size and only using OpenACC for larger ones
 - or do we even need the smaller grids?