

Case study: the parallel Himeno code

Alistair Hart


Cray Exascale Research Initiative Europe





Contents

- **A parallel code is a scalar code with data transfers**
 - We have looked at how to port a scalar code
 - Here we look at the parallel version of the same code
- **The new feature is the data transfer between processors**
 - which means local data transfers between CPU and GPU
- **This talk looks at the extra things we need to consider**
 - First at a conceptual level
 - the OpenACC part
 - Then some specific points for two different comm models
 - MPI
 - Fortran coarrays



The Himeno Benchmark

- **Parallel 3D Poisson equation solver**
 - Iterative loop evaluating 19-point stencil
 - Memory intensive, memory bandwidth bound
- **Fortran, C, MPI and OpenMP implementations available from <http://acc.riken.jp/2444.htm>**
- **Fortran Coarray (CAF) version developed**
 - ~600 lines of Fortran
 - Fully ported to accelerator using 27 directive pairs
- **Strong scaling benchmark**
 - XL configuration: 1024 x 512 x 512 global volume
 - Expect halo exchanges to become significant as we scale
 - Async. GPU data transfers and kernel launches to help avoid this

Overall program structure

- Like scalar case:

- initmt() initialises data
- jacobi(nn,gosa)
 - does nn iterations
 - stencil update to data
- called twice:
 - once for calibration
 - once for measurement

- differences:

- initcomm() routine
 - sets up processor grid

```

PROGRAM himeno
  CALL initcomm      ! Set up processor grid
  CALL initmt        ! Initialise local matrices

  cpu0 = gettime() ! Wraps SYSTEM_CLOCK
  CALL jacobi(3,gosa)
  cpu1 = gettime()
  cpu = cpu1 - cpu0

  ! nn = INT(ttargt/(cpu/3.0)) ! Fixed runtime
  nn = 1000                ! Hardwired for testing

  cpu0 = gettime()
  CALL jacobi(nn,gosa)
  cpu1 = gettime()
  cpu = cpu1 - cpu0

  xmflops2 = flop*1.0d-6/cpu*nn

  PRINT *, ' Loop executed ',nn,' times'
  PRINT *, ' Gosa : ',gosa
  PRINT *, ' MFLOPS: ',xmflops2,' time(s): ',cpu
END PROGRAM himeno

```



The distributed jacobi routine

- iteration loop:
 - fixed tripcount
- jacobi kernel:
 - new pressure array `wrk2`
 - local residual `wgosa`
- halo exchange
 - between neighbours
 - uses send, receive buffers
- Allreduce
 - global residual `gosa`
 - `wgosa` summed over PEs
- second kernel:
 - `p` array updated from `wrk2`

```
DO loop = 1, nn

    compute Jacobi: wrk2, wgosa

    pack halos from wrk2 into send bufs

    exchange halos with neighbour PEs

    Allreduce to sum wgosa across Pes

    copy back wrk2 into p

    unpack halos into p from recv bufs

ENDDO
```

The Jacobi computational kernel (serial)

- The stencil is applied to pressure array **p**
- Updated pressure values are saved to temporary array **wrk2**
- Local control value **wgosa** is computed

```

DO K=2,kmax-1
  DO J=2,jmax-1
    DO I=2,imax-1
      s0=a(I,J,K,1)* p(I+1,J, K )
      +a(I,J,K,2)* p(I, J+1,K ) &
      +a(I,J,K,3)* p(I, J, K+1) &
      +b(I,J,K,1)*(p(I+1,J+1,K )-p(I+1,J-1,K ) &
      -p(I-1,J+1,K )+p(I-1,J-1,K )) &
      +b(I,J,K,2)*(p(I, J+1,K+1)-p(I, J-1,K+1) &
      -p(I, J+1,K-1)+p(I, J-1,K-1)) &
      +b(I,J,K,3)*(p(I+1,J, K+1)-p(I-1,J, K+1) &
      -p(I+1,J, K-1)+p(I-1,J, K-1)) &
      +c(I,J,K,1)* p(I-1,J, K ) &
      +c(I,J,K,2)* p(I, J-1,K ) &
      +c(I,J,K,3)* p(I, J, K-1) &
      + wrk1(I,J,K)

      s = (s0*a(I,J,K,4)-p(I,J,K))*bnd(I,J,K)
      wgosa = wgosa + ss*ss
      wrk2(I,J,K) = p(I,J,K) + omega * ss
    ENDDO
  ENDDO
ENDDO

```

fwd n.n. n.n. bwd n.n.

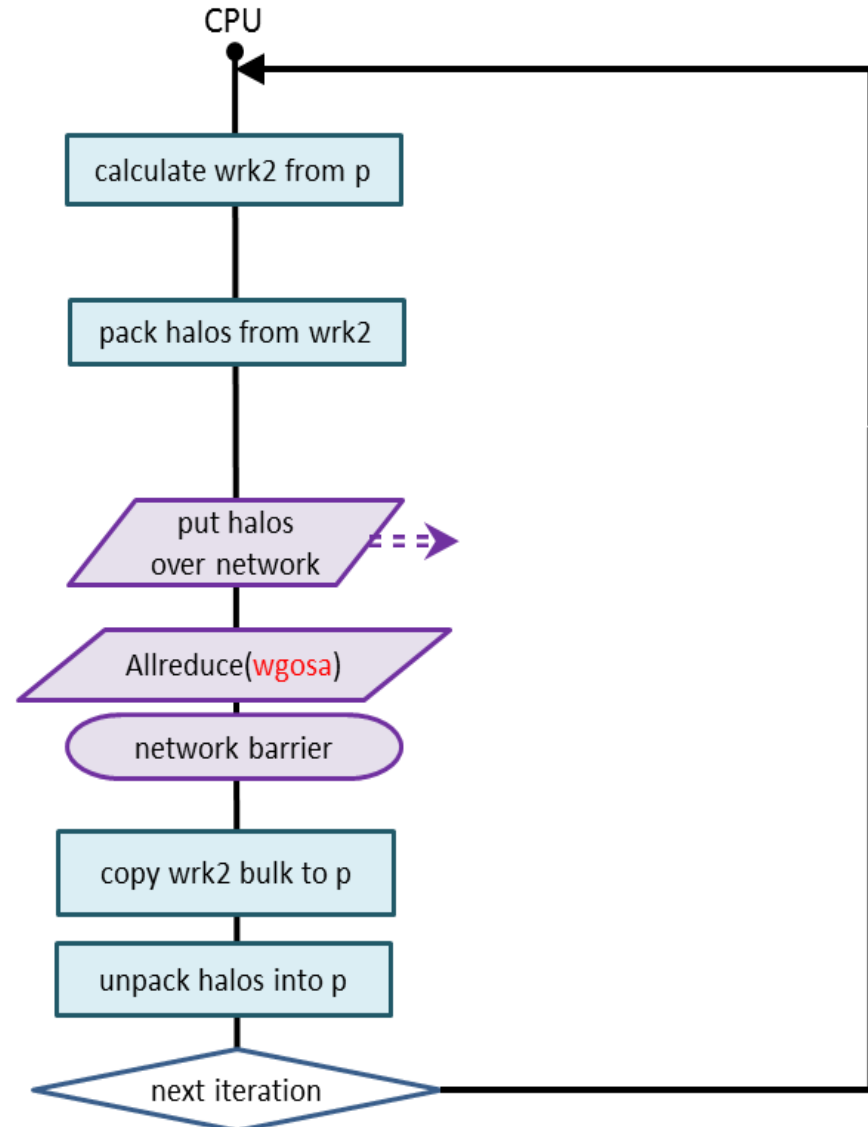
Distributed jacobi routine as a flowchart

- tasks reorganised

- take advantage of overlap
 - just network at this point
- still some freedom:
 - barrier can move relative to
 - Allreduce
 - wrk2→bulk

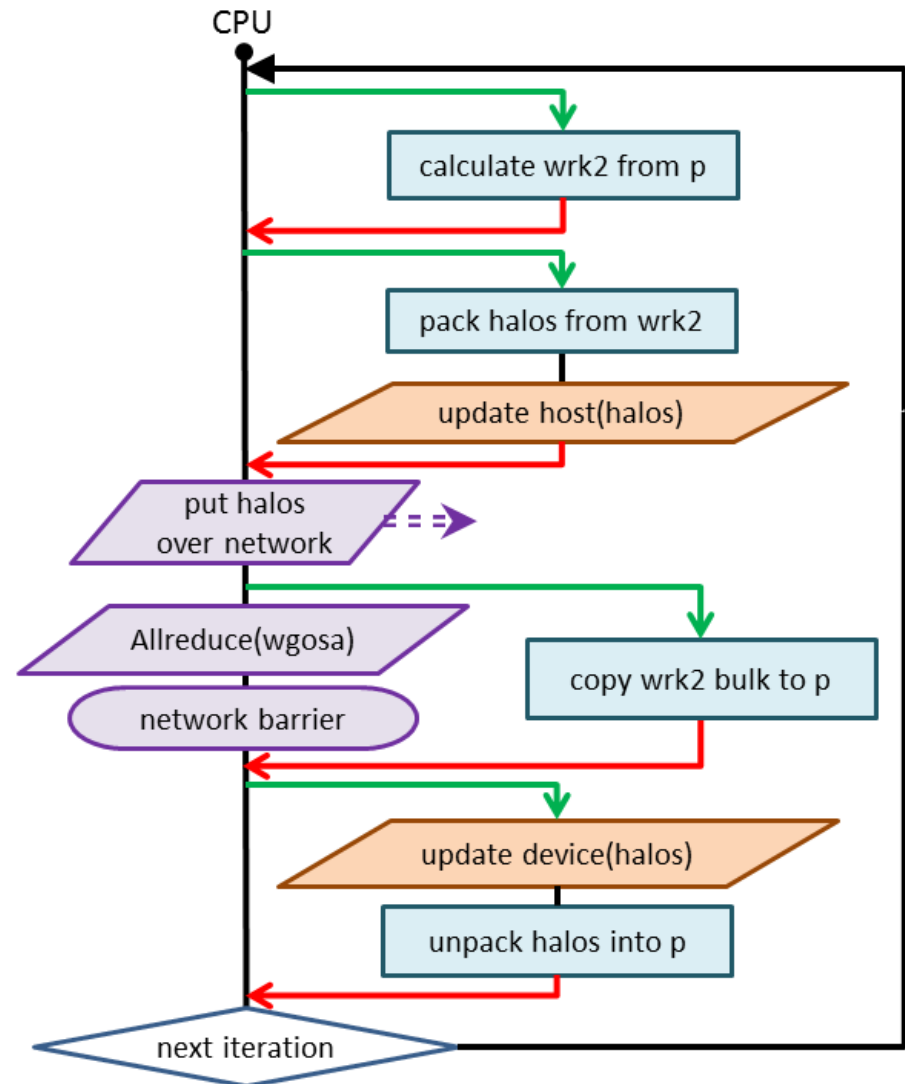
- Which is best order?

- depends on:
 - hardware
 - comms library
 - kernel details
 - local problem size



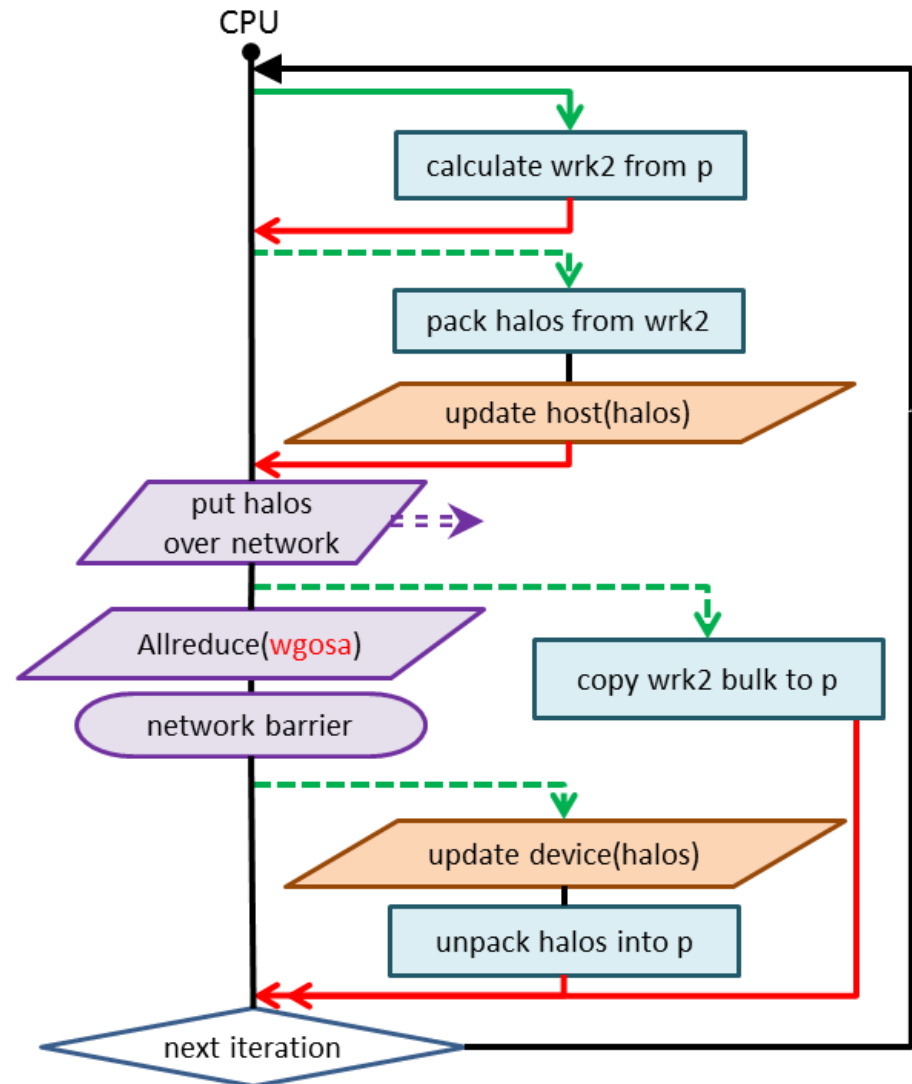
First OpenACC port

- **Loopnest kernels**
 - stencil
 - pack halos
 - wrk2→bulk
 - unpack halos
- **acc updates**
- **Some additional overlap**
 - GPU kernels asynchronous
 - wrk2→bulk
 - can overlap comms



Asynchronicity

- **async handle**
 - allows task overlap on GPU
- **Dotted lines are async**
 - halo pack
 - can overlap x, y, z directions
 - halo unpack
 - can overlap x, y, z directions
 - also overlap with wrk2→bulk
- **Red lines are sync points**





Packing and transferring send buffers

- **Use async clause**

- separate handles for overlap
- one per direction
 - three in total
 - integers, e.g. 1, 2, 3
- what about six?
 - one each for up and down
 - not as efficient
 - but this is a tuning option
- global wait
 - all 3 streams completed

```
!$acc parallel loop async(xstrm)
DO k = 2,kmax-1
  DO j = 2,jmax-1
    sendbuffx_dn(j,k)=wrk2(2,j,k)
    sendbuffx_up(j,k)=wrk2(imax-1,j,k)
  ENDDO
ENDDO
!$acc end parallel loop

!$acc update host &
!$acc (sendbuffx_dn,sendbuffx_up) &
!$acc async(xstrm)

! same for y with async(ystrm)
! same for z with async(zstrm)

!$acc wait
<send the 6 buffers>
```

Packing and transferring send buffers (better)

• The biggest bottleneck

- PCIe link serialises:
 - one buffer at a time
- So:
 - as soon as one direction done
 - send these buffers over network
- Ordering:
 - Can we guarantee **x** ready first?
 - No, but likely
- OpenACC does not have **waitany** directive
 - so go with most likely ordering

```
!$acc parallel loop async(xstrm)
<pack the two x buffers>
!$acc end parallel loop

!$acc update host &
!$acc  (sendbuffx_dn, sendbuffx_up) &
!$acc  async(xstrm)

! same for y with async(ystrm)
! same for z with async(zstrm)

!$acc wait(xstrm)
<send the two x buffers>

!$acc wait(ystrm)
<send the two y buffers>

!$acc wait(zstrm)
<send the two z buffers>
```

Transferring and unpacking recv buffers

- **Copying wrk2 into p**
 - very quick
 - but can overlap network
- **When messages complete**
 - copy and unpack recv buffers
 - three parallel streams
- **wait**
 - ensures all 4 streams complete

```

!$acc parallel loop async(bstrm)
<wrk2 -> p>
!$acc end parallel loop

<network barrier>

!$acc update device &
!$acc (recvbuffx_dn,recvbuffx_up) &
!$acc async(xstrm)

!$acc parallel loop async(xstrm)
<unpack the two x buffers>
!$acc end parallel loop

! same for y with async(ystrm)
! same for z with async(zstrm)

!$acc wait

<end iteration loop>

```



Data region considerations

- **Twelve extra buffers**
 - 3 physical dimensions: x, y, z
 - 2 directions of transfer: up, down
 - Separate buffers for send and receive
- **These buffers need to be added to the data region(s)**
 - if just in `jacobi()` data region
 - they should be **create**
 - allocated each time `jacobi()` is called, no data transfer (except explicit updates)
 - or in outer data region
 - **create** in main data region
 - **present** in `jacobi()` data region
 - if `jacobi()` routine was called a lot, better to allocate once in parent

Communications models

- **Two domain decompositions strategies**
 - Both message passing:
 - MPI
 - Fortran coarrays (CAF)
- **A lot of similarities**
 - a few subtle differences



MPI

- **Use nonblocking MPI calls**
 - **MPI_IRECV**
 - post these receives early; first thing in iteration loop
 - **MPI_ISEND**
 - call these in 3 sets (**x**, **y**, **z**) of two (up, down)
 - after each of **async** streams **xstrm**, **ystrm**, **zstrm** separately completes
 - network barrier:
 - **MPI_WAITALL**(12,send_handles,recv_handles)
 - ensures all buffers arrive (and send buffers can be reused)
 - But, PCIe transfer of recv buffers is a bottleneck
 - better to start transfer of messages to GPU as soon as they arrive
 - don't know which order 6 message will arrive in
 - **MPI_WAITANY**(6,recv_handles)
 - As each arrives, start async stream of: update, then unpack kernel
 - After all recvs completed:
 - **MPI_WAITALL**(6, send_handles)
- **Could investigate use of nonblocking collective for gosa**
 - MPI 3.0



Fortran coarrays

- **Introduced in Fortran2008**

- previously known as Co-array Fortran (CAF) language extension

- **PGAS feature**

- single-sided communication model

- **Simplifies the code**

- e.g. to send a buffer:
 - `recvbuffx_up(:,:)[myx-1,myy,myz] = sendbuffx_dn(:,:)`
- send buffer is a (local) array, but a (global) **co-scalar**
- recv buffer is a (local) array and a (global) **coarray**
- no acknowledgement, receipt etc. for message being sent
- RDMA put: remote "image" doesn't know when data arrives
- user has to handle all synchronisation points
- gives great potential for overlapping comms with compute etc.



CAF and Himeno

- **"put" operations are more efficient than "get"**
 - i.e. put remote **coarray** on the left of the "="
 - `recvbuffx_up(:,:)[myx-1,myy,myz] = sendbuffx_dn(:,:)`
- **But, Fortran standard says that:**
 - Not only should this guarantee `z = x`:
 - `y = x`
 - `z = y`
 - But also, so should this:
 - `y[remote] = x`
 - `z = y[remote]`
- So compiler forced to add a lot of (usually unnecessary) sync points
- To avoid these, need CCE directive above CAF statements
 - `!$pgas defer_sync`

Using CAF for parallel Himeno with OpenACC



- **send buffers**

- **co-scalars**, no problem using them in OpenACC directives

- **recv buffers**

- **co-arrays**, forbidden from being in OpenACC directives
 - to prevent race conditions
- so we need to:
 - copy local recv (**coarray**) buffer into temporary (**co-scalar**) array on CPU
 - use the temporary array in the update, parallel loop kernels
- additional CPU buffer copy potentially affects performance/scalability

- **data region**

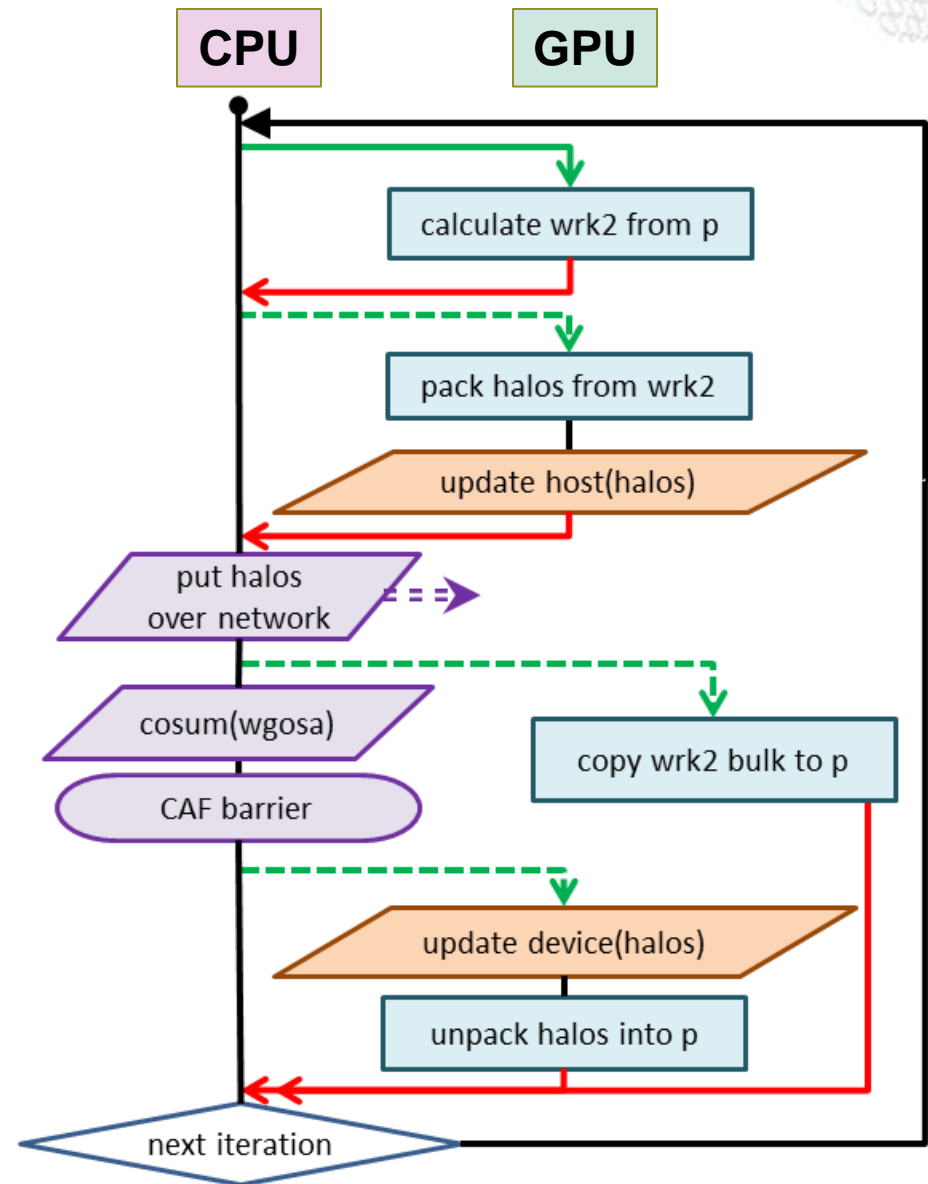
- send buffers and temporary buffers (all **co-scalars**) go in **data** region
 - using **create** or **present** clause

CAF synchronisation

- **Network synchronisation**

- needed to guarantee the halo buffers have arrived
- `sync all`
- but we are also globally summing the residual
 - scalar `wgosa` needs to be a `co-scalar` (as used in OpenACC `reduction`)
 - copy `wgosa` to local part of `coarray cgosa`
 - `gosa = cosum(cgosa)`
- `cosum` partially synchronises the images
 - `sync all` would now be overkill
 - `sync memory` instead completes the synchronisation
- this guarantees all messages have completed
 - no equivalent of `MPI_WAITANY` to relieve PCIe recv buffer congestion

Final CAF implementation



OpenACC / CAF version

- **Total number of lines in the original Himeno MPI-Fortran code:** 629
- **Total number lines in the modified version with coarrays and accelerator directives:** 554
 - don't need MPI_CART_CREATE and the like
- **Total number of accelerator directives:** 27
 - plus 18 "end" directives



Benchmarking the code

- **Cray XK6 configuration:**

- Single AMD IL-16 2.1GHz nodes, 16 cores per node
- Nvidia Tesla X2090 GPU, 1 GPU per node
- Running with 1 PE (GPU) per node
- Himeno case XL needs at least 16 XK6 nodes
- Testing blocking and asynchronous GPU implementations


- **Cray XE6 configuration:**

- Dual AMD IL-16 2.1 GHz nodes, 32 cores per node
- Running on fully packed nodes: all cores used
- Depending on the number of nodes, 1-4 OpenMP threads per PE are used

- **All comparisons are for strong scaling on case XL**

Himeno performance

- XK6 GPU is about 1.6x faster than XE6
- OpenACC async implementation is ~ 8% faster than OpenACC blocking





Himeno code breakdown

- **Host/GPU transfers take more time than the halo exchange (network)**
 - this code would benefit from an efficient direct GPU-GPU communication
- **On 128 nodes, ~55% of the time is spent in the GPU compute kernel**

