

# Practical: a sample code

Alistair Hart  
Cray Exascale Research Initiative Europe



# Aims

- The aim of this practical is to examine, compile and run a simple, pre-prepared OpenACC code

# The sample code

- **The sample code**
  - Designed to demonstrate functionality
  - not interested in performance at this stage
- **Implements the simple example from the lectures**
  - A 3d array **a** is initialised
  - It's values are doubled and stored in a new array **b**
  - A checksum is calculated and compared with the expected result
- **These are implemented as 3 OpenACC kernels**
- **There are three versions of the code**
  - Version 00 has all 3 kernels in same main program.
    - There is no attempt to keep data on the GPU between the kernels.
  - Version 01 uses a data region to avoid data sloshing.
  - Version 02 has more complicated calltree
    - calls a subroutine that contains an OpenACC kernel.
    - This kernel also contains a function call.



# Code versions and building them

- **There are versions for 4 different programming models**
  - C or Fortran, with
  - static or dynamic allocation of arrays
    - N.B. there is no version00 for dynamic arrays with C (see note in version01)
  - source filename based on these, e.g. `first_example_Fstatic_v00.f90`
- **Get your environment right**
  - make sure you have the right PrgEnv loaded (cray or pgi)
  - make sure you have loaded the correct compiler version module
  - make sure you have loaded module **craype-accel-nvidia20**
- **Build the code**
  - PrgEnv-cray:
    - `ftn -hlist=a <Fortran source file>`
    - `cc -hlist=a <C source file>`
  - PrgEnv-pgi:
    - `ftn -Minfo=all <Fortran source file>`
    - `cc -Minfo=all <C source file>`



# Automation

- You can do it all by hand if you wish, or use automation
  - There's nothing magic being done here
- **Automated building:**
  - can just type: `make VERSION=[00|01|02] [F|C][static|dynamic]`
    - Makefile will echo commands it uses to build the code
    - automatically detects which PrgEnv you are using (uses PE\_ENV env. var.)
      - remember to type "`make clean`" if you switch PrgEnv modules
- **Automated building and job submission**
  - type: `bash build_submit.bash MYPE TARGET VERSION`
    - MYPE should be `cray` or `pgi`
    - TARGET should be `Fstatic` or `Fdynamic` or `Cstatic` or `Cdynamic`
    - VERSION should be `00` or `01` or `02`
  - This will:
    - load the correct modules using script `../XK_setup.bash`
    - build the code using the **Makefile**
    - create directory: `/lus/scratch/$USER/Cray_OpenACC_training/Practical1/TARGET_VERSION_DATE_TIME`
    - write and submit a PBS jobscript
  - You can then `cd` to this directory and look at the output

# What to check

- **Check correctness**

- Did the code compile correctly?
- Did the job execute?
- Was the answer correct?

- **Next, understand what did the compiler did**

- examine and understand the compiler feedback
  - CCE: open the .lst file
  - PGI: read the output to stdout
- did it compile for the accelerator?
- what data did it plan to move and when?
- how were the loop iterations scheduled?



# What actually ran?

- **Did we actually run on the accelerator?**

- We can ask the runtime for some feedback
- cd to run directory, edit jobscript, uncomment appropriate line
  - CCE: set **CRAY\_ACC\_DEBUG** to 1 (least detailed) to 3 (most detailed)
  - PGI: set **ACC\_NOTIFY**
- Resubmit the job: **qsub <jobscript name>**
- Examine commentary (in the log file) and make sure you understand it

- **Profiling the code**

- A quick way of profiling is to use the Nvidia compute profiler
  - CCE and PGI compile to PTX (as does nvcc), so this will work for all
- Edit the jobscript and uncomment the profiling line
- Resubmit the job
- Examine the profile (in file **cuda\_profile\_0.log**)
  - Can change location with env. var. **COMPUTE\_PROFILE\_LOG**
- This is a "blow-by-blow" account
  - Larger codes need a more aggregated report
  - We will cover profiling in more detail later

## Further work

- Choose a target and repeat this for all three versions
- Start with the Cray compiler
- Then either:
  - repeat for a different programming model target, or
  - try the PGI compiler





# Getting the examples

- **On raven:**

- change to a directory where you want to work
  - either in your home directory or under `/lus/scratch/$USER`
- type: `tar zxvf ~tr99/cray_OpenACC_training.tgz`
- This creates a new directory `./Cray_OpenACC_training`
  - please note the file `LICENCE.txt`
- The codes for Practical 1 are in:
  - `Cray_OpenACC_training/Practical1`
  - There is a README file that summarises these slides
- `~tr99/Slides/*` contains PDFs of the slides supporting the practicals